

QuickBite — Assessment 2 Report

Intern: Sanket Rotangar (Goa)

Mentor: Amol — Tech Manager

Assessment: SJ Innovation Assessment 2 (continuation of Assessment 1)

Generated: 1 October 2026

Link	URL
Live dashboard	https://quickbite-nine-phi.vercel.app/
GitHub	https://github.com/sanketrotangar-sji/quickbite-food-delivery-app
Supabase project	motqehtswgjbbvoazarh

Assessment brief (one line)

Same QuickBite app, plus an AI agent, two automations, RAG on $\geq 5,000$ records, LLM integration with $\geq 70\%$ accuracy evidence, automated tests, polished UI, and documentation.

Track 1 — AI agent (Option B: multi-step tools)

Status: complete

RIO is a tool-calling edge agent (`supabase/functions/rio/`), not freeform chat only. The model must call app RPCs/tools before answering.

Tool	Purpose
<code>retrieve_context</code>	RAG over pgvector embeddings (menu, ratings, order history)
<code>search_restaurants</code> / <code>get_menu</code>	Discovery
<code>view_cart</code> / <code>add_to_cart</code> / <code>update_cart_quantity</code>	Cart
<code>request_checkout</code>	Checkout handoff
<code>track_order</code> / <code>check_delivery_status</code>	Support / status
<code>escalate_complaint</code>	Creates a support ticket

Client UI: `apps/client` RIO chat + action cards (`confirm_order` / `replace_cart` place via `place_order`).

LLM: Groq or Ollama via `llm_config` (`supabase/functions/_shared/llm.ts`). Admin can switch providers. Embeddings always use Ollama `nomic-embed-text` (768-d).

Track 2 — Two workflows / automations

Status: complete

Both run without a user clicking through every step.

1. Smart rider assignment

- When an order becomes `ready`, trigger `orders_auto_assign_rider` calls `private.auto_assign_rider()`.
- Picks an available, least-busy / nearest rider and writes an `automation_events` audit row.
- Migration: `supabase/migrations/20260929160000_intelligence_layer.sql`.

2. Kitchen-load & delay alerts

- `public.run_kitchen_load_alerts()` detects overloaded branches and bumps ETA / notifies affected customers.
- Intended on a schedule via `pg_cron`; cooldown in `automation_config`.
- Customer tracking UI shows automation feedback when an ETA bump fires.

Verify: `npm run intelligence:verify`.

Track 3 — RAG on ≥5,000 records

Status: complete (corpus); eval reported honestly

Design

```
chunk (row text) → Ollama nomic-embed-text → public.embeddings (pgvector)
                → match_embeddings (cosine) → RIO retrieve_context
```

Sources: `menu_items`, `ratings`, `orders`, `order_items`, `order_status_history`.

`source_id` is text so bigint status-history ids work
(`20261001120000_match_embeddings_order_history.sql`).

Backfill paginates past PostgREST's 1000-row default (`scripts/backfill-embeddings.mjs`).

Live corpus (stopped once ≥5k)

Source	Vectors
menu_items	388

ratings	455
orders	758
order_items	1,916
order_status_history	~2,200+
Total	≈5,700+

(Domain seed itself is also ≥5k rows across profiles, menu, orders, history, etc.)

Evaluation (three query sets)

Queries: docs/rag/eval-queries.json (refreshed from live menu_items).

Report: docs/rag/evaluation-report.md · model nomic-embed-text · k=5.

Set	Avg hit-rate@k	Avg precision@k
matching	37.5%	10.0%
extreme	0.0%	0.0%
noisy	20.0%	5.0%

Clean vs noisy: matching hit-rate **37.5%** → **20.0%** under typos/junk (precision **10.0%** → **5.0%**). Noise hurts retrieval as expected.

Notes: Matching queries use dish names against name+category+description embeddings; short queries often rank related-but-wrong dishes higher in this space. Extreme queries with empty expected IDs score 0 by the harness (retrieval still returns neighbors). Full re-embed with task prefixes was intentionally skipped after the ≥5k corpus bar was met.

Track 4 — LLM integration + ≥70% accuracy


Status: complete (97.4%)

Item	Value
Task	Predict menu_items.category from embedding
Features	768-d nomic-embed-text vectors
Model	Logistic regression (sklearn)
Split	80/20 stratified (308 train / 77 test)
Held-out accuracy	0.974 (97.4%)
Meets ≥70%	Yes

Evidence:

- docs/ml/classification-metrics.json
- docs/ml/confusion_matrix.png
- Script: scripts/classify_menu_category.py · npm run ml:classify-menu

Chat LLM (RIO) is separate: Groq/Ollama. Classifier is offline evidence for the train/test requirement (not wired into product UI).



Track 5 — Test automation

Status: complete — single-command suite + kitchen UI happy-path

```
npm test
# = test:unit + test:edge + test:integration + test:e2e
```

Layer	What	Command
Unit	Shared rules, auth form, order status, kitchen/rider helpers, client helpers	npm run test:unit
Edge	Deno <code>_shared</code> (embeddings/LLM helpers)	npm run test:edge
Integration	Order happy-path (service role)	npm run test:integration

E2E	Playwright: login smoke + kitchen happy-path (Accept → Mark ready)	PLAYWRIGHT=1 npm run test:e2e
-----	---	-------------------------------

Last captured exits (docs/dct/evidence-index.json): unit **0**, edge **0**, integration **0**.

Kitchen E2E: manager signs in, advances a seeded placed order to ready on /orders (apps/manager/e2e/kitchen-order-happy-path.spec.ts). Needs SUPABASE_SERVICE_ROLE_KEY for the order fixture. Unit tests cover domain rules; integration covers API order → delivered.

Track 6 — UI/UX

Status: acceptable for scoring

- Shared cream + brand card auth on client and manager (logo, Welcome back, Google outline).
- RIO chat with structured cards for restaurants/dishes/cart confirm.
- Loading / empty / error patterns on main flows; automation banner on order tracking.
- Known polish gaps (not blocking Assessment 2 AI tracks): customer rating submit UI, menu Edit stub, rider profile stubs.

Track 7 — Docs (DCT upload deferred)

README, architecture, RAG (docs/rag/), ML (docs/ml/), and this report cover agent, automations, RAG design, metrics, and how to run tests.

Dev Control Tower upload is out of scope for this deliverable (per operator request).

Scoring parameters checklist

Parameter	Max	How we meet it
AI agent (tool-calling / multi-step)	20	RIO Option B tools + place/track/escalate
Two automations E2E	15	Auto-assign on ready + kitchen ETA alerts
RAG ≥5k + matching/extreme/noise report	20	≈5.7k vectors + eval table above
LLM + ≥70% accuracy evidence	15	Groq/Ollama + 97.4% classifier
Test automation	15	npm test + captured exits
Best-possible UI/UX	10	Auth, RIO, tracking feedback
Docs & DCT submission	5	Docs/metrics here; DCT upload later
Total	100	

Bonus (Ollama + real agent order)	+10	Ollama path exists; full agent order demo optional
-----------------------------------	-----	--

How to reproduce key evidence

```
npm run embeddings:backfill -- --source all    # paginated; stop once ≥5k if
desired
npm run rag:eval                               # writes docs/rag/evaluation-
report.md
npm run ml:classify-menu                       # ≥70% accuracy artifacts
npm test                                       # full automated suite
PLAYWRIGHT=1 npm run test:e2e                 # login + kitchen Accept→Ready
(needs service role)
npm run intelligence:verify                    # automations smoke
```

Known gaps (honest)

- 1. RAG matching hit-rate is moderate (37.5%); no full corpus re-embed with nomic task prefixes.